

LAB Assignment No. 5

Topic: Artificial Neural Network Model

Question 1

Logic Gates with Neural Network. Implement a feed-forward neural network to learn the AND gate.

- Inputs: (0,0), (0,1), (1,0), (1,1)
- Output: 0, 0, 0, 1

Tasks:

1. Create dataset using NumPy or pandas.
2. Build a neural network with one hidden layer using TensorFlow/Keras or PyTorch.
3. Train it and show accuracy.
4. Compare model predictions with actual outputs

Code:

```
import numpy as np
from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import Dense
X = np.array([[0,0], [0,1], [1,0], [1,1]])
y = np.array([[0], [0], [0], [1]])

model = Sequential()
model.add(Dense(4, input_dim=2, activation='relu')) # Hidden layer with 4 neurons
model.add(Dense(1, activation='sigmoid')) # Output layer

model.compile(optimizer='adam', loss='binary_crossentropy', metrics=['accuracy'])
model.fit(X, y, epochs=500, verbose=0)
loss, accuracy = model.evaluate(X, y, verbose=0)
print(f"Model Accuracy: {accuracy*100:.2f}%")
predictions = model.predict(X)
predicted_classes = (predictions > 0.5).astype(int)
```

Output:

```
Model Accuracy: 100.00%
1/1 ━━━━━━━━ 0s 102ms/step

Predictions vs Actual:
Input: [0 0] Predicted: 0 Actual: 0
Input: [0 1] Predicted: 0 Actual: 0
Input: [1 0] Predicted: 0 Actual: 0
Input: [1 1] Predicted: 1 Actual: 1
```

Question 2

Create a dataset $y = x^2 + \text{noise}$ for x in range [-3,3]. Regression Task with Neural Network

Tasks:

1. Generate 100 samples.
2. Build a neural network to predict y from x .
3. Plot actual vs. predicted results.
4. Discuss how increasing hidden neurons changes results.


Code:

```
import numpy as np
import matplotlib.pyplot as plt
from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import Dense
np.random.seed(42)
X = np.linspace(-3, 3, 100)
y = X**2 + np.random.normal(0, 0.5, X.shape) # Add Gaussian noise
X = X.reshape(-1, 1)
y = y.reshape(-1, 1)
model = Sequential()
model.add(Dense(10, input_dim=1, activation='relu')) # Hidden layer (10 neurons)
model.add(Dense(1)) # Output layer (regression)

model.compile(optimizer='adam', loss='mse')
model.fit(X, y, epochs=200, verbose=0)

y_pred = model.predict(X)
plt.figure(figsize=(8,5))
plt.scatter(X, y, label='Actual', color='blue')
plt.plot(X, y_pred, label='Predicted', color='red', linewidth=2)
plt.title("Regression:  $y = x^2 + \text{noise}$ ")
plt.xlabel("x")
plt.ylabel("y")
plt.legend()
plt.show()
```

Output:



Question 3:

Use the XOR gate and train networks with different activation functions (sigmoid, tanh, ReLU).

- Compare accuracy, loss, and convergence speed.
- Plot and discuss results.

Code:

```
import numpy as np
import matplotlib.pyplot as plt
from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import Dense
from tensorflow.keras.optimizers import Adam
X = np.array([[0,0], [0,1], [1,0], [1,1]])
y = np.array([[0], [1], [1], [0]])
def train_xor_model(activation, epochs=500):
    model = Sequential([
        Dense(4, input_dim=2, activation=activation),
        Dense(1, activation='sigmoid')
    ])
```

```

model.compile(optimizer=Adam(learning_rate=0.1),
              loss='binary_crossentropy',
              metrics=['accuracy'])
history = model.fit(X, y, epochs=epochs, verbose=0)
loss, acc = model.evaluate(X, y, verbose=0)
return model, history, loss, acc
activations = ['sigmoid', 'tanh', 'relu']
results = {}
for act in activations:
    model, history, loss, acc = train_xor_model(act)
    results[act] = {'loss': loss, 'acc': acc, 'history': history}
    print(f'{act.upper()} → Accuracy: {acc*100:.2f}%, Final Loss: {loss:.4f}')
plt.figure(figsize=(8,5))
for act in activations:
    plt.plot(results[act]['history'].history['loss'], label=f'{act} loss')
plt.title("Training Loss vs Epochs for Different Activations")
plt.xlabel("Epochs")
plt.ylabel("Loss")
plt.legend()
plt.grid(True)
plt.show()


```

Output:

```

SIGMOID → Accuracy: 100.00%, Final Loss: 0.0031
TANH → Accuracy: 100.00%, Final Loss: 0.0002
RELU → Accuracy: 100.00%, Final Loss: 0.0007

```



Question 4


Binary Classification using Neural Network

Objective: Build a neural network to classify whether a tumor is malignant or benign using the **Breast Cancer dataset**.

Code:

```
import numpy as np
from sklearn.datasets import load_breast_cancer
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler
from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import Dense
data = load_breast_cancer()
X = data.data
y = data.target # 0 = malignant, 1 = benign
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3, random_state=42)
scaler = StandardScaler()
X_train = scaler.fit_transform(X_train)
X_test = scaler.transform(X_test)
model = Sequential([
    Dense(16, input_dim=X.shape[1], activation='relu'), # hidden layer 1
    Dense(8, activation='relu'), # hidden layer 2
    Dense(1, activation='sigmoid')])
model.compile(optimizer='adam', loss='binary_crossentropy', metrics=['accuracy'])
history = model.fit(X_train, y_train, validation_data=(X_test, y_test), epochs=100, batch_size=16,
verbose=0)
loss, accuracy = model.evaluate(X_test, y_test, verbose=0)
print(f"✅ Test Accuracy: {accuracy*100:.2f}%")
print(f"Loss: {loss:.4f}")
import matplotlib.pyplot as plt
plt.figure(figsize=(12,5))
plt.subplot(1,2,1)
plt.plot(history.history['accuracy'], label='Train Accuracy')
plt.plot(history.history['val_accuracy'], label='Test Accuracy')
plt.title("Model Accuracy")
plt.xlabel("Epochs")
plt.ylabel("Accuracy")
plt.legend()
plt.subplot(1,2,2)
plt.plot(history.history['loss'], label='Train Loss')
plt.plot(history.history['val_loss'], label='Test Loss')
plt.title("Model Loss")
plt.xlabel("Epochs")
plt.ylabel("Loss")
plt.legend()
plt.show()
sample = X_test[0].reshape(1, -1)
prediction = model.predict(sample)
print(f"Predicted Class: {'Benign' if prediction[0][0] > 0.5 else 'Malignant'}")
```

Output:



Question 5

Multi-Class Classification on Iris Dataset

Objective: Train a neural network to classify flower species (Setosa, Versicolor, Virginica).

Code:

```
import numpy as np
import matplotlib.pyplot as plt
from sklearn.datasets import load_iris
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler, LabelEncoder
from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import Dense
from tensorflow.keras.utils import to_categorical
iris = load_iris()
X = iris.data          # Features: sepal/petal length & width
y = iris.target         # Target: 0 = setosa, 1 = versicolor, 2 = virginica
y_encoded = to_categorical(y)
X_train, X_test, y_train, y_test = train_test_split(
    X, y_encoded, test_size=0.3, random_state=42)
```

```

scaler = StandardScaler()
X_train = scaler.fit_transform(X_train)
X_test = scaler.transform(X_test)

model = Sequential([
    Dense(8, input_dim=4, activation='relu'), # hidden layer 1
    Dense(6, activation='relu'), # hidden layer 2
    Dense(3, activation='softmax') # output layer (3 neurons = 3 classes)
])
model.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])

history = model.fit(X_train, y_train, validation_data=(X_test, y_test),
                     epochs=100, batch_size=8, verbose=0)

loss, accuracy = model.evaluate(X_test, y_test, verbose=0)
print(f" ✅ Test Accuracy: {accuracy*100:.2f}%")
print(f"Loss: {loss:.4f}")

plt.figure(figsize=(12,5))
plt.subplot(1,2,1)
plt.plot(history.history['accuracy'], label='Train Accuracy')
plt.plot(history.history['val_accuracy'], label='Test Accuracy')
plt.title('Model Accuracy')
plt.xlabel('Epochs')
plt.ylabel('Accuracy')
plt.legend()


plt.subplot(1,2,2)
plt.plot(history.history['loss'], label='Train Loss')
plt.plot(history.history['val_loss'], label='Test Loss')
plt.title('Model Loss')
plt.xlabel('Epochs')
plt.ylabel('Loss')
plt.legend()
plt.show()

sample = np.array([[5.1, 3.5, 1.4, 0.2]]) # Example: Setosa
sample_scaled = scaler.transform(sample)
prediction = model.predict(sample_scaled)
predicted_class = np.argmax(prediction)
print(f"Predicted Species: {iris.target_names[predicted_class]}")

```

Output:

 Test Accuracy: 97.78%
 Loss: 0.0953



```
1/1 ━━━━━━━━ 0s 107ms/step
Predicted Species: setosa
```

Question 6

Regression Problem (House Price Prediction)

Objective: Predict house prices using the **California Housing dataset**.

```
import numpy as np
import matplotlib.pyplot as plt
from sklearn.datasets import fetch_california_housing
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler
from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import Dense

# 1 Load the dataset
data = fetch_california_housing()
X = data.data
y = data.target # Median house value (in 100,000s)

# 2 Split into training and testing sets
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)

scaler = StandardScaler()
X_train = scaler.fit_transform(X_train)
X_test = scaler.transform(X_test)
```

```

model = Sequential([
    Dense(64, input_dim=X.shape[1], activation='relu'), # Hidden layer 1
    Dense(32, activation='relu'), # Hidden layer 2
    Dense(1) # Output layer (regression)])

model.compile(optimizer='adam', loss='mse', metrics=['mae']) # MSE: Mean Squared Error, MAE: Mean Absolute Error

history = model.fit(X_train, y_train, validation_data=(X_test, y_test),
                     epochs=100, batch_size=32, verbose=0)
loss, mae = model.evaluate(X_test, y_test, verbose=0)
print(f" ✅ Test MAE (Mean Absolute Error): {mae:.3f}")
print(f"Test MSE: {loss:.3f}")

plt.figure(figsize=(12,5))
plt.subplot(1,2,1)
plt.plot(history.history['loss'], label='Train Loss')
plt.plot(history.history['val_loss'], label='Test Loss')
plt.title('Model Loss (MSE)')
plt.xlabel('Epochs')
plt.ylabel('Loss')
plt.legend()
plt.subplot(1,2,2)
plt.plot(history.history['mae'], label='Train MAE')
plt.plot(history.history['val_mae'], label='Test MAE')
plt.title('Model Mean Absolute Error')
plt.xlabel('Epochs')
plt.ylabel('MAE')
plt.legend()
plt.show()
sample = X_test[0].reshape(1, -1)
predicted_price = model.predict(sample)[0][0]
print(f"Predicted Price: ${predicted_price * 100000:.2f}")


```

Output:

```

✅ Test MAE (Mean Absolute Error): 0.345
Test MSE: 0.262

```



Question 7

Neural Network with Dropout Regularization

Objective: Prevent overfitting using **Dropout** layers on the **MNIST digit dataset**.

Code:

```
import numpy as np
import matplotlib.pyplot as plt
from tensorflow.keras.datasets import mnist
from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import Dense, Dropout, Flatten
from tensorflow.keras.utils import to_categorical
(X_train, y_train), (X_test, y_test) = mnist.load_data()
X_train = X_train.astype('float32') / 255.0
X_test = X_test.astype('float32') / 255.0
X_train = X_train.reshape(-1, 28*28)
X_test = X_test.reshape(-1, 28*28)

y_train = to_categorical(y_train, 10)
y_test = to_categorical(y_test, 10)

model = Sequential([
    Dense(512, input_dim=784, activation='relu'),
    Dropout(0.3),          # 30% of neurons randomly dropped
    Dense(256, activation='relu'),
    Dropout(0.3),          # another Dropout layer
    Dense(10, activation='softmax') # output layer (10 classes)])
model.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])
```

```

history = model.fit(
    X_train, y_train,
    validation_data=(X_test, y_test),
    epochs=20,
    batch_size=128,
    verbose=1)
loss, accuracy = model.evaluate(X_test, y_test, verbose=0)
print(f"✓ Test Accuracy: {accuracy*100:.2f}%")
print(f"Test Loss: {loss:.4f}")
plt.figure(figsize=(12,5))
plt.subplot(1,2,1)
plt.plot(history.history['accuracy'], label='Train Accuracy')
plt.plot(history.history['val_accuracy'], label='Test Accuracy')
plt.title('Accuracy with Dropout')
plt.xlabel('Epochs')
plt.ylabel('Accuracy')
plt.legend()
plt.subplot(1,2,2)
plt.plot(history.history['loss'], label='Train Loss')
plt.plot(history.history['val_loss'], label='Test Loss')
plt.title('Loss with Dropout')
plt.xlabel('Epochs')
plt.ylabel('Loss')
plt.legend()
plt.show()

```

Output:

```

11490434/11490434 0s 0us/step
Epoch 1/20
469/469 14s 22ms/step - accuracy: 0.8350 - loss: 0.5230 - val_accuracy: 0.9668 - val_loss: 0.1107
Epoch 2/20
469/469 9s 19ms/step - accuracy: 0.9608 - loss: 0.1280 - val_accuracy: 0.9717 - val_loss: 0.0834
Epoch 3/20
469/469 9s 19ms/step - accuracy: 0.9720 - loss: 0.0910 - val_accuracy: 0.9755 - val_loss: 0.0757
Epoch 4/20
469/469 10s 21ms/step - accuracy: 0.9777 - loss: 0.0695 - val_accuracy: 0.9779 - val_loss: 0.0641
Epoch 5/20
469/469 10s 20ms/step - accuracy: 0.9820 - loss: 0.0575 - val_accuracy: 0.9809 - val_loss: 0.0631
Epoch 6/20
469/469 8s 18ms/step - accuracy: 0.9843 - loss: 0.0492 - val_accuracy: 0.9805 - val_loss: 0.0640
Epoch 7/20
469/469 10s 20ms/step - accuracy: 0.9861 - loss: 0.0426 - val_accuracy: 0.9821 - val_loss: 0.0584
Epoch 8/20
469/469 10s 21ms/step - accuracy: 0.9866 - loss: 0.0411 - val_accuracy: 0.9822 - val_loss: 0.0627
Epoch 9/20
469/469 9s 19ms/step - accuracy: 0.9878 - loss: 0.0372 - val_accuracy: 0.9815 - val_loss: 0.0615
Epoch 10/20
469/469 10s 19ms/step - accuracy: 0.9892 - loss: 0.0329 - val_accuracy: 0.9828 - val_loss: 0.0615
Epoch 11/20
469/469 11s 21ms/step - accuracy: 0.9901 - loss: 0.0309 - val_accuracy: 0.9853 - val_loss: 0.0592
Epoch 12/20
...
Epoch 20/20
469/469 11s 21ms/step - accuracy: 0.9938 - loss: 0.0192 - val_accuracy: 0.9848 - val_loss: 0.0687
✓ Test Accuracy: 98.48%
Test Loss: 0.0687

```

