

LAB Assignment No 7

Recurrent Neural Network (RNN)

LAB Task 1:

Next Word Prediction using RNN

Objective: Learn how RNNs can predict the next word in a sentence.

Dataset: Any small text corpus — e.g., *Shakespeare.txt* or *Wikipedia sample*.

Tasks:

1. Load and clean the text data.
2. Tokenize and convert text into sequences.
3. Build a simple **RNN model** using keras.layers.SimpleRNN.
4. Train it to predict the next word given previous 3–5 words.
5. Test by entering a custom text prompt and predict the next word.

```
import numpy as np
import tensorflow as tf
from tensorflow.keras.preprocessing.text import Tokenizer
from tensorflow.keras.preprocessing.sequence import pad_sequences
from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import Embedding, SimpleRNN, Dense
text = """
The sun is shining bright
The sun is hot today
The weather is sunny
The sun is very bright
"""
text = text.lower().strip()
tokenizer = Tokenizer()
tokenizer.fit_on_texts([text])
```

```

total_words = len(tokenizer.word_index) + 1

input_sequences = []

for line in text.split("\n"):

    token_list = tokenizer.texts_to_sequences([line])[0]

    for i in range(1, len(token_list)):

        input_sequences.append(token_list[:i+1])

max_seq_len = max(len(seq) for seq in input_sequences)

input_sequences = pad_sequences(input_sequences,
                                 maxlen=max_seq_len,
                                 padding='pre')

X = input_sequences[:, :-1]

y = input_sequences[:, -1]

y = tf.keras.utils.to_categorical(y, num_classes=total_words)

model = Sequential()

model.add(Embedding(total_words, 50, input_length=max_seq_len-1))

model.add(SimpleRNN(100))

model.add(Dense(total_words, activation='softmax'))

model.compile(loss='categorical_crossentropy',
               optimizer='adam',
               metrics=['accuracy'])

model.summary()

model.fit(X, y, epochs=200, verbose=1)

def predict_next_word(model, tokenizer, text, max_seq_len):

    text = text.lower()

    sequence = tokenizer.texts_to_sequences([text])[0]

    sequence = pad_sequences([sequence],
                            maxlen=max_seq_len-1,
                            padding='pre')

```

```

prediction = model.predict(sequence, verbose=0)

predicted_index = np.argmax(prediction)

for word, index in tokenizer.word_index.items():

    if index == predicted_index:

        return word

input_text = "the sun is"

predicted_word = predict_next_word(model, tokenizer, input_text, max_seq_len)

print(f"Input: '{input_text}'")

print(f"Predicted Next Word: '{predicted_word}'")

```

Output:

```

Input: 'the sun is'
Predicted Next Word: 'very'

```

LAB Task 2:

Stock Price Prediction using RNN

Objective: Predict future stock prices using time series data.

Dataset: Use Google Stock Price dataset (from Kaggle or Yahoo Finance).

Tasks:

1. Import dataset and normalize values.
2. Prepare time-step sequences (e.g., 60 previous days → next day price).
3. Build and train an **RNN model** using SimpleRNN layers.
4. Evaluate predictions vs actual prices (plot graph).

```
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
from sklearn.preprocessing import MinMaxScaler
from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import SimpleRNN, Dense
data = {
    "Date": [
        "2023-01-02", "2023-01-03", "2023-01-04", "2023-01-05", "2023-01-06",
        "2023-01-09", "2023-01-10", "2023-01-11", "2023-01-12", "2023-01-13",
        "2023-01-16", "2023-01-17", "2023-01-18", "2023-01-19", "2023-01-20",
        "2023-01-23", "2023-01-24", "2023-01-25", "2023-01-26", "2023-01-27",
        "2023-01-30", "2023-01-31", "2023-02-01", "2023-02-02", "2023-02-03",
        "2023-02-06", "2023-02-07", "2023-02-08", "2023-02-09", "2023-02-10",
        "2023-02-13", "2023-02-14", "2023-02-15", "2023-02-16", "2023-02-17",
        "2023-02-20", "2023-02-21", "2023-02-22", "2023-02-23", "2023-02-24",
        "2023-02-27", "2023-02-28", "2023-03-01", "2023-03-02", "2023-03-03",
        "2023-03-06", "2023-03-07", "2023-03-08", "2023-03-09", "2023-03-10"
    ],
    "Close": [
        2685.5, 2690.2, 2688.75, 2701.3, 2710.4,
        2705.6, 2718.9, 2725.1, 2719.8, 2732.5,
        2740.2, 2735.9, 2748.6, 2755.3, 2762.1,
        2812.4, 2820.9, 2828.6, 2835.2, 2842.9,
        2838.5, 2846.1, 2852.8, 2860.4, 2868.9,
        2865.3, 2872.7, 2879.5, 2886.2, 2893.8,
        2890.1, 2898.6, 2905.2, 2912.8, 2920.4,
        2916.9, 2925.5, 2932.1, 2938.7, 2945.3
    ]
}
```

```
2786.9,2795.4,2802.1,2810.3,2818.7,  
2812.4,2820.9,2828.6,2835.2,2842.9,  
2838.5,2846.1,2852.8,2860.4,2868.9,  
2865.3,2872.7,2879.5,2886.2,2893.8,  
2890.1,2898.6,2905.2,2912.8,2920.4,  
2916.9,2925.5,2932.1,2938.7,2945.3  
]  
}  
  
dataset = pd.DataFrame(data)  
  
prices = dataset[['Close']].values  
  
scaler = MinMaxScaler(feature_range=(0, 1))  
  
prices_scaled = scaler.fit_transform(prices)  
  
X = []  
  
y = []  
  
for i in range(10, len(prices_scaled)):  
    X.append(prices_scaled[i-10:i, 0])  
    y.append(prices_scaled[i, 0])  
  
X = np.array(X)  
  
y = np.array(y)  
  
X = np.reshape(X, (X.shape[0], X.shape[1], 1))  
  
model = Sequential()  
  
model.add(SimpleRNN(50, activation='tanh', input_shape=(10, 1)))  
  
model.add(Dense(1))  
  
model.compile(optimizer='adam', loss='mean_squared_error')  
  
model.summary()  
  
model.fit(X, y, epochs=20, batch_size=8)  
  
predicted_prices = model.predict(X)  
  
predicted_prices = scaler.inverse_transform(predicted_prices)  
  
real_prices = scaler.inverse_transform(y.reshape(-1, 1))
```

```

plt.figure(figsize=(10, 6))

plt.plot(real_prices, color='blue', label='Actual Stock Price')

plt.plot(predicted_prices, color='red', label='Predicted Stock Price')

plt.title('Stock Price Prediction using SimpleRNN')

plt.xlabel('Time')

plt.ylabel('Stock Price')

plt.legend()

plt.show()

```

Output:

```

Model: "sequential"


```

Layer (type)	Output Shape	Param #
simple_rnn (SimpleRNN)	(None, 50)	2,600
dense (Dense)	(None, 1)	51

```

Total params: 2,651 (10.36 KB)
Trainable params: 2,651 (10.36 KB)
Non-trainable params: 0 (0.00 B)

```



LAB Task 3:

Sentiment Analysis using RNN

Objective: Classify movie reviews as positive or negative using RNN.

Dataset: *IMDb Movie Reviews* dataset (available in Keras).

Tasks:

1. Load dataset and preprocess text (tokenize and pad sequences).
2. Build RNN with Embedding + SimpleRNN layers.
3. Train for binary classification (positive/negative).
4. Evaluate accuracy on test data.

```
import numpy as np

from tensorflow.keras.datasets import imdb

from tensorflow.keras.models import Sequential

from tensorflow.keras.layers import Embedding, SimpleRNN, Dense

from tensorflow.keras.preprocessing.sequence import pad_sequences

vocab_size = 10000

(x_train, y_train), (x_test, y_test) = imdb.load_data(num_words=vocab_size)

max_len = 200

x_train = pad_sequences(x_train, maxlen=max_len)

x_test = pad_sequences(x_test, maxlen=max_len)

model = Sequential([
    Embedding(input_dim=vocab_size, output_dim=128, input_length=max_len),
    SimpleRNN(64, activation='tanh'),
    Dense(1, activation='sigmoid')])

model.compile(
    loss='binary_crossentropy',
    optimizer='adam',
    metrics=['accuracy'])

model.fit(
    x_train,
```

```

y_train,
epochs=5,
batch_size=64,
validation_split=0.2)

loss, accuracy = model.evaluate(x_test, y_test)

print(f"Test Accuracy: {accuracy * 100:.2f}%")

word_index = imdb.get_word_index()

def encode_review(text):

    encoded = []

    for word in text.lower().split():

        index = word_index.get(word)

        if index is not None and index < vocab_size:

            encoded.append(index + 3)

    return pad_sequences([encoded], maxlen=max_len)

custom_text = "This movie was amazing and full of emotions"

encoded_review = encode_review(custom_text)

prediction = model.predict(encoded_review)

if prediction[0][0] > 0.5:

    print("Sentiment: Positive Review")

else:

    print("Sentiment: Negative Review")

```

Output:

```

1641221/1641221 ━━━━━━━━ 0s
1/1 ━━━━━━━━ 0s 179ms/step
Sentiment: Negative Review

```

LAB Task 4:

Weather Forecasting using RNN

Objective: Predict future temperature based on previous days' readings.

Dataset: Daily temperature dataset (e.g., “Jena Climate Dataset” from TensorFlow).


Tasks:

1. Load and visualize temperature over time.
2. Prepare input-output sequences for time series prediction.
3. Build an RNN to predict next day's temperature.
4. Plot actual vs predicted temperature.

```
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
from sklearn.preprocessing import MinMaxScaler
from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import SimpleRNN, Dense
data = pd.read_csv("daily_temperature_dataset.csv")
temperature = data["Temperature (degC)"].values
plt.figure()
plt.plot(temperature[:2000])
plt.title("Temperature Over Time")
plt.xlabel("Time")
plt.ylabel("Temperature (°C)")
plt.show()
scaler = MinMaxScaler()
temperature_scaled = scaler.fit_transform(temperature.reshape(-1, 1))
def create_sequences(data, seq_length):
    X, y = [], []
    for i in range(len(data) - seq_length):
        X.append(data[i:i+seq_length])
        y.append(data[i+seq_length])
    return np.array(X), np.array(y)
```

```
for i in range(len(data) - seq_length):  
    X.append(data[i:i + seq_length])  
    y.append(data[i + seq_length])  
  
return np.array(X), np.array(y)  
  
sequence_length = 30  
  
X, y = create_sequences(temperature_scaled, sequence_length)  
  
split = int(0.8 * len(X))  
  
X_train, X_test = X[:split], X[split:]  
y_train, y_test = y[:split], y[split:]  
  
model = Sequential([  
    SimpleRNN(50, activation='tanh', input_shape=(sequence_length, 1)),  
    Dense(1)])  
  
model.compile(optimizer='adam', loss='mse')  
  
model.fit(X_train, y_train,  
          epochs=10,  
          batch_size=32,  
          validation_split=0.2)  
  
predicted = model.predict(X_test)  
  
predicted_temp = scaler.inverse_transform(predicted)  
actual_temp = scaler.inverse_transform(y_test)  
  
plt.figure()  
  
plt.plot(actual_temp, label="Actual Temperature")  
plt.plot(predicted_temp, label="Predicted Temperature")  
plt.title("Actual vs Predicted Temperature")  
plt.xlabel("Days")  
plt.ylabel("Temperature (°C)")  
plt.legend()  
plt.show()
```

Output:



LAB Task 5:

Music Note Generation using RNN

Objective: Generate new music sequences using RNN.

Dataset: *MIDI music dataset* (short sequences or melodies).

Tasks:

1. Convert MIDI data into integer-encoded notes.
2. Train an RNN on note sequences (input: previous notes → output: next note).

```
import numpy as np

from music21 import note, chord, stream

from tensorflow.keras.models import Sequential

from tensorflow.keras.layers import SimpleRNN, Dense

notes = [

    "C4","D4","E4","F4","G4","A4","B4","C5",

    "C5","B4","A4","G4","F4","E4","D4","C4",

    "C4.E4.G4","D4.F4.A4","E4.G4.B4","F4.A4.C5"

]

unique_notes = sorted(set(notes))

note_to_int = {n:i for i,n in enumerate(unique_notes)}

int_to_note = {i:n for i,n in enumerate(unique_notes)}

sequence_length = 4

X, y = [], []

for i in range(len(notes)-sequence_length):

    seq_in = notes[i:i+sequence_length]

    seq_out = notes[i+sequence_length]

    X.append([note_to_int[n] for n in seq_in])

    y.append(note_to_int[seq_out])
```

```
X = np.reshape(X, (len(X), sequence_length, 1))

X = X / float(len(unique_notes))

y = np.array(y)

model = Sequential()

model.add(SimpleRNN(128, input_shape=(X.shape[1], X.shape[2])))

model.add(Dense(len(unique_notes), activation='softmax'))

model.compile(loss='sparse_categorical_crossentropy', optimizer='adam')

model.fit(X, y, epochs=100, batch_size=1, verbose=0)

start = np.random.randint(0, len(X)-1)

pattern = X[start]

generated_notes = []

for _ in range(10):

    prediction = model.predict(pattern.reshape(1, sequence_length, 1), verbose=0)

    index = np.argmax(prediction)

    result = int_to_note[index]

    generated_notes.append(result)

    next_input = index / float(len(unique_notes))

    pattern = np.append(pattern, [[next_input]], axis=0)

    pattern = pattern[1:]

print("Generated Notes:", generated_notes)

output_notes = []

offset = 0

for pattern in generated_notes:

    if '.' in pattern:

        chord_notes = pattern.split('.')

        chord_objects = [note.Note(n) for n in chord_notes]

        new_chord = chord.Chord(chord_objects)

        new_chord.offset = offset
```

```
output_notes.append(new_chord)

else:

    new_note = note.Note(pattern)
    new_note.offset = offset
    output_notes.append(new_note)
    offset += 0.5

midi_stream = stream.Stream(output_notes)

midi_stream.write('midi', fp='generated_music_demo.mid')
```

Output:

```
Generated Notes: ['F4', 'E4', 'D4', 'C4', 'C4.E4.G4', 'D4.F4.A4', 'E4.G4.B4', 'F4.A4.C5', 'G4', 'A4']
MIDI file generated: generated_music_demo.mid
```