

# LAB No 12

## Implementation of Reinforcement Learning Part 1

### Experiment: CartPole Environment using Gymnasium & Pygame

---

#### 🎯 Lab Objectives

After completing this lab, students will be able to:

- Understand the **Reinforcement Learning interaction loop**
- Use **Gymnasium environments**
- Visualize agent behavior using **Pygame**
- Interpret **states, actions, rewards, and episodes**
- Modify and analyze RL environment parameters

```
import gymnasium as gym
import pygame

env = gym.make("CartPole-v1", render_mode="human")

font = None

for episode in range(1, 20):
    score = 0
    state, info = env.reset()
    done = False

    while not done:
        action = env.action_space.sample()
        state, reward, terminated, truncated, info = env.step(action)
        done = terminated or truncated
        score += reward


        if font is None:
            pygame.font.init()
            font = pygame.font.SysFont("Arial", 24)

        surface = pygame.display.get_surface()
```

```
text = font.render(f"Score: {int(score)}", True, (255, 0, 0))
surface.blit(text, (10, 10))
pygame.display.update()

print(f"Episode {episode} Score: {score}")

env.close()
pygame.quit()
```



## Lab Questions (Conceptual Understanding)

### Q1.

What is **Reinforcement Learning**? Identify the **agent**, **environment**, **state**, **action**, and **reward** in the given code.

**Answer:**

RL is a type of machine learning where an agent learns by **taking actions in an environment** and getting **rewards**. The goal is to maximize total rewards over time.

In the CartPole code:

Component	Description in the Code
<b>Agent</b>	The part of the code that chooses actions: <code>action = env.action_space.sample()</code>
<b>Environment</b>	The CartPole simulation created with: <code>env = gym.make("CartPole-v1", render_mode="human")</code>
<b>State</b>	The current situation of the system returned by <code>env.reset()</code> or <code>env.step(action)</code> It includes: [cart_position, cart_velocity, pole_angle, pole_angular_velocity]
<b>Action</b>	What the agent does at each step: 0 = push cart left, 1 = push cart right
<b>Reward</b>	The feedback the agent gets: <code>reward = 1</code> for every step the pole stays upright

---

### Q2.

Explain the purpose of the following line:

```
env = gym.make("CartPole-v1", render_mode="human")
```

**Answer:**

This line **sets up the environment** and makes it **visible to the user** so you can watch the agent act in real time.

---

### Q3.

What does env.reset() return? Why are two values returned?

**Answer:**

**What it does:**

- env.reset() **resets the environment** to the starting state for a new episode.
- It returns **two values**:
  1. **state** → The initial observation of the environment (CartPole variables: cart position, cart velocity, pole angle, pole angular velocity)
  2. **info** → Additional information from the environment (usually empty or extra metadata; not used in basic experiments)

**Why two values:**

- Gymnasium separates **the important observation (state)** from **optional metadata (info)**.
  - This allows the code to use the state for decision-making while still having access to extra info if needed.
- 

### Q4.

Explain the difference between: Terminated and truncated

**Answer:**

Term	Meaning in Gymnasium / CartPole
<b>terminated</b>	The episode ended because the <b>goal was reached</b> or <b>failure occurred</b> . In CartPole: the pole fell too far or cart moved out of bounds.
<b>truncated</b>	The episode ended because it <b>reached the maximum allowed steps</b> . This is not due to failure, just a <b>time limit</b> .

---

## **Q5.**

What is the role of the variable score? How is it calculated?

**Answer:**

**Role:**

- score keeps track of the **total reward** the agent receives during an episode.
- It measures **how well the agent is performing**—higher score means the pole stayed upright longer.

**How it is calculated:**

score += reward

- At each step, the environment gives a **reward** (in CartPole, reward = 1 per step).
  - The code **adds this reward to score** until the episode ends.
  - At the end of the episode, score represents the **total steps the pole stayed balanced**.
- 

## **Q6.**

Why is action = env.action\_space.sample() used?

Is this an intelligent agent? Justify your answer.

**Answer:**

action = env.action\_space.sample() is used to **choose a random action** at each step.

**This is not an intelligent agent** because it **does not learn** from rewards or past experience; it acts **randomly**.

---

## **Q7.**

Explain how **Pygame** is used to display the score on the screen.

**Answer:**

**Pygame** is used to **draw the score on the simulation window**.

**Steps in the code:**

1. Initialize Pygame font: `pygame.font.SysFont("Arial", 24)`
  2. Create a surface (the window) using `pygame.display.get_surface()`
  3. Render the score as text: `font.render(f"Score: {int(score)}", True, (255,0,0))`
  4. Draw it on the window at a position: `surface.blit(text, (10, 10))`
  5. Update the display: `pygame.display.update()`
- 

## **Q8.**

What happens if the `pygame.display.update()` line is removed?

**Answer:**

- If `pygame.display.update()` is removed, the **score will not appear or refresh** on the screen.
- The **Pygame** window **won't show changes**, so the score text won't be visible while the simulation runs.

## Lab Tasks (Hands-on Practice)

### ◆ Task 1: Modify Number of Episodes

Change the number of episodes from **20 to 50** and observe:

- How the score varies across episodes
- Whether performance improves or remains random

Answer:


- The score **changes a lot from episode to episode.**
  - Some episodes have **low scores** (around 9–15).
  - Some episodes have **higher scores** (around 40–58).
  - There is **no fixed pattern** in the scores
- 

### ◆ Task 2: Display Episode Number on Screen

Modify the code to show:


Episode: X | Score: Y

on the CartPole window.



## Task 3: Change Text Color and Position

- Change score text color from **red to green**
- Display it at position **(200, 20)**



---

#### ◆ Task 4: Print Maximum Score

After all episodes finish:

- Store all episode scores
- Print the **maximum score achieved**

```
Maximum Score: 59.0
(venv) PS D:\AI> []
```


---

#### ◆ Task 5: Slow Down the Environment

Insert a small delay using:

```
pygame.time.delay(20)
```

Observe the effect on visualization.




## ◆ Task 6: Replace CartPole with MountainCar

Change the environment to:

```
env = gym.make("MountainCar-v0", render_mode="human")
```

Compare:

- Reward behavior
- Episode termination condition



## Task 7: Identify State Variables

Print the state vector and answer:

- How many state variables are there?
- What does each variable represent?


```
Episode 19 Score: -200.0
```

---

## ◆ Task 8 (Advanced): Rule-Based Action

Replace random action with:

```
if state[2] > 0:  
    action = 1  
else:  
    action = 0
```



### Observation Table (For Students)

Episode	Score	Remarks	Print
1			
2			
...			
20			

## Experiment: MountainCar Environment using Gymnasium & Pygame

### Lab Objectives

After completing this lab, students will be able to:

- Understand the **working of a continuous control RL environment**
- Analyze **delayed reward problems**
- Use **Gymnasium MountainCar-v0**
- Visualize agent behavior and rewards using **Pygame**
- Compare MountainCar with CartPole environment

**Provided Code:**

```
import gymnasium as gym
import pygame

env = gym.make("MountainCar-v0", render_mode="human")

font = None
best_score = -float('inf')

# We only need a few episodes to prove it works with a better policy
NUM_EPISODES = 5

for episode in range(1, NUM_EPISODES + 1):
    state, info = env.reset()
    done = False
    score = 0

    while not done:
        # Task 7/8: Advanced Rule-Based Action
        # state[1] is velocity. If velocity is moving right (>0), push right (2).
        # If moving left (<0), push left (0). This builds momentum rapidly.
        if state[1] > 0:
            action = 2
        else:
            action = 0

        state, reward, terminated, truncated, info = env.step(action)
        done = terminated or truncated
        score += reward


        if font is None:
            pygame.font.init()
            font = pygame.font.SysFont("Arial", 24)

        surface = pygame.display.get_surface()
        text = font.render(f"Episode: {episode} Score: {int(score)}", True, (0, 0, 255))
        surface.blit(text, (200, 20))

        # Reduced delay for faster execution
        pygame.time.delay(5)
        pygame.display.update()

    print(f"Episode {episode} Score: {score}")
```

```
if score > best_score:  
    best_score = score  
  
env.close()  
pygame.quit()  
  
print(f"\nOptimization Results:")  
print(f"Best Score Achieved: {best_score}")
```



## Lab Questions (Conceptual Understanding)

### Q1.

What is **Reinforcement Learning**? Identify the **agent**, **environment**, **state**, **action**, and **reward** in the MountainCar code.

**Answer:**

**Reinforcement Learning (RL)** is a learning method where an agent learns by interacting with an environment and receiving rewards.

**In MountainCar code:**

- **Agent:** The car controller (our program)
  - **Environment:** MountainCar-v0
  - **State:** [position, velocity]
  - **Action:** Push left (0), no push (1), push right (2)
  - **Reward:** -1 at every step until goal is reached
- 

### Q2.

Explain the purpose of the following statement:

```
env = gym.make("MountainCar-v0", render_mode="human")
```

**Answer:**

- Creates the MountainCar environment
  - render\_mode="human" displays the environment visually on the screen
- 

### Q3.

What are the **state variables** in MountainCar-v0? What does each state represent?

**Answer:**

- **Position (state[0])**  
→ Horizontal position of the car on the hill

- **Velocity (state[1])**  
→ Speed and direction of the car's movement
- 

#### Q4.

Describe the **action space** of MountainCar-v0. How many actions are available and what do they mean?

**Answer:**

Action	Meaning
0	Push car left
1	No push
2	Push car right

---

#### Q5.

Explain the reward mechanism in MountainCar-v0.

Why does the agent receive a **negative reward** at each step?

**Answer:**

- The agent receives **-1 reward at every step**
- The goal is to **reach the hilltop in fewer steps**

**Reason for negative reward:**

To encourage the agent to reach the goal as quickly as possible.

---

#### Q6.

What is the difference between:

Terminated and truncated in this environment?

**Answer:**

- **terminated:** Episode ends because the goal is reached
  - **truncated:** Episode ends because maximum step limit is reached
-

## **Q7.**

Why does the agent fail to reach the goal when using `action_space.sample()`?

### **Answer:**

- Random actions do not build momentum
  - The car cannot climb the hill without coordinated left-right movement
  - Therefore, the agent fails to reach the goal
- 

## **Q8.**

Explain the role of **momentum** in solving the MountainCar problem.

### **Answer:**

- The car must first move **away from the goal** to gain speed
  - Momentum helps the car climb the steep hill
  - Without momentum, the car cannot reach the top
-